

Crude Oil Price Predictor 2020

Christopher Page

8/12/2020

Introduction/Overview/Executive Summary

Oil is an important commodity. Some of our planet's inhabitants produce it. Many more consume it. All are affected by it. Because of its wide-ranging effects, oil commands collective attention enabled by analysis. One of the key metrics to take into account in the analysis of the oil industry is the daily closing price of crude oil futures traded on a global basis in such venues as the New York Mercantile Exchange.

Methods for predicting that price may entail studying the quantitative and qualitative factors at play inside the oil industry as well as in such adjacent industries as transportation and manufacturing. They may also entail studying the often complex diplomatic, military, economic, cultural, and environmental developments that help explain why crude oil closing prices rise and fall as they do over time.

This project proposes a simpler method, one that predicts crude oil closing prices based on time and the closing prices of complementary (e.g., gasoline) and competing (e.g., platinum) commodities. The premise is that there is much to be learned by studying the behaviors of commodity traders engaged in making daily investment decisions based on running evaluations of risks and rewards.

That premise informed the development of algorithms with the potential to predict crude oil closing prices with a viable level of accuracy, quantified as the Root Means Square Error (RMSE). Using a 25,190-row data set constructed with publicly available information downloaded for free from <https://www.nasdaq.com/>, the author developed and then evaluated fifteen such algorithms.

Evaluation led to the selection of one algorithm that generated the lowest RMSE when applied to a test set. That algorithm used a Ranger model which took into account six key predictors:

- (1) Year
- (2) Month
- (3) Closing Price of Heating Oil
- (4) Closing Price of Gasoline
- (5) Closing Price of Platinum
- (6) Closing Price of Soybeans

The RMSE in question was 2.93, a figure equating to a mere 12% of the 23.51 generated by an algorithm that only took into account the average closing price across the period of observation.

Method/Analysis

```
## Loading required package: caret
```

```
## Loading required package: lattice
```

```
## Loading required package: ggplot2
```

```

## Loading required package: caTools

## Warning: package 'caTools' was built under R version 4.0.2

## Loading required package: dplyr

##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union

## Loading required package: dslabs

## Loading required package: forcats

## Loading required package: foreach

## Loading required package: gam

## Warning: package 'gam' was built under R version 4.0.2

## Loading required package: splines

## Loaded gam 1.20

## Loading required package: ggrepel

## Loading required package: ggthemes

## Loading required package: lubridate

##
## Attaching package: 'lubridate'

## The following objects are masked from 'package:base':
##
##   date, intersect, setdiff, union

## Loading required package: purrr

##
## Attaching package: 'purrr'

```

```

## The following objects are masked from 'package:foreach':
##
##   accumulate, when

## The following object is masked from 'package:caret':
##
##   lift

## Loading required package: randomForest

## randomForest 4.6-14

## Type rfNews() to see new features/changes/bug fixes.

##
## Attaching package: 'randomForest'

## The following object is masked from 'package:dplyr':
##
##   combine

## The following object is masked from 'package:ggplot2':
##
##   margin

## Loading required package: tidyverse

## -- Attaching packages -----

## v tibble 3.0.1      v readr 1.3.1
## v tidyr  1.1.0      v stringr 1.4.0

## -- Conflicts -----
## x purrr::accumulate() masks foreach::accumulate()
## x lubridate::as.difftime() masks base::as.difftime()
## x randomForest::combine() masks dplyr::combine()
## x lubridate::date() masks base::date()
## x dplyr::filter() masks stats::filter()
## x lubridate::intersect() masks base::intersect()
## x dplyr::lag() masks stats::lag()
## x purrr::lift() masks caret::lift()
## x randomForest::margin() masks ggplot2::margin()
## x lubridate::setdiff() masks base::setdiff()
## x lubridate::union() masks base::union()
## x purrr::when() masks foreach::when()

## Parsed with column specification:
## cols(
##   Date = col_character(),
##   Commodity = col_character(),
##   Sector = col_character(),
##   'Closing Price' = col_double()
## )

```

```

## Joining, by = "date"
## Joining, by = "date"
## Joining, by = "date"
## Joining, by = "date"
## Joining, by = "date"
## Joining, by = "date"
## Joining, by = "date"
## Joining, by = "date"
## Joining, by = "date"

## [1] "Crude Oil Commodity Prices began the August 2010 to July 2020 period at"

## [1] 78.95

## [1] "They ended that ten-year period at"

## [1] 40.34

## [1] "During that time, they averaged"

## [1] 70.13307

## [1] "Running from a minimum of"

## [1] -37.25

## [1] "To a maximum of"

## [1] 113.45

## [1] "With a standard deviation of"


## [1] 23.52209

### Second, this script plots the 2010-2020 timeline of crude oil daily closing prices
### Then, it saves the plot in a .png file entitled "crude_oil_price_history"

crude_oil_plot <- crude_oil %>%
  ggplot(aes(date, closing_price)) +
  geom_point(aes(color = commodity)) +
  xlab("Date") +
  ylab("Closing Price") +
  ggtitle("Crude Oil Price History") +
  theme_economist() +
  theme(legend.position="top",
        legend.title = element_blank(),
        legend.box = "horizontal" ,
        legend.text=element_text(size=8.5)) +
  guides(col = guide_legend(nrow = 1))


crude_oil_plot

```




```
ggsave("fig/crude_oil_price_history.png")
```

```
## Saving 6.5 x 4.5 in image
```




Saving 6.5 x 4.5 in image




```
## Saving 6.5 x 4.5 in image
```

```
crude_oil_average_by_quarter_of_the_year_plot
```




```
ggsave("fig/crude_oil_average_by_quarter_of_the_year.png")
```

```
## Saving 6.5 x 4.5 in image
```


Saving 6.5 x 4.5 in image



```
## Saving 6.5 x 4.5 in image
```


```
crude_oil_vs_other_energy_plot
```

Crude Oil versus Other Energy Sector Commodities




```
ggsave("fig/crude_oil_in_comparison_to_three_other_energy_sector_commodities.png")
```

```
## Saving 6.5 x 4.5 in image
```




```
## Saving 6.5 x 4.5 in image
```

```
crude_oil_heating_oil_scatterplot
```




```
ggsave("fig/scatterplot_crude_oil_versus_heating_oil.png")
```


```
## Saving 6.5 x 4.5 in image
```




```
## Saving 6.5 x 4.5 in image
```



Saving 6.5 x 4.5 in image

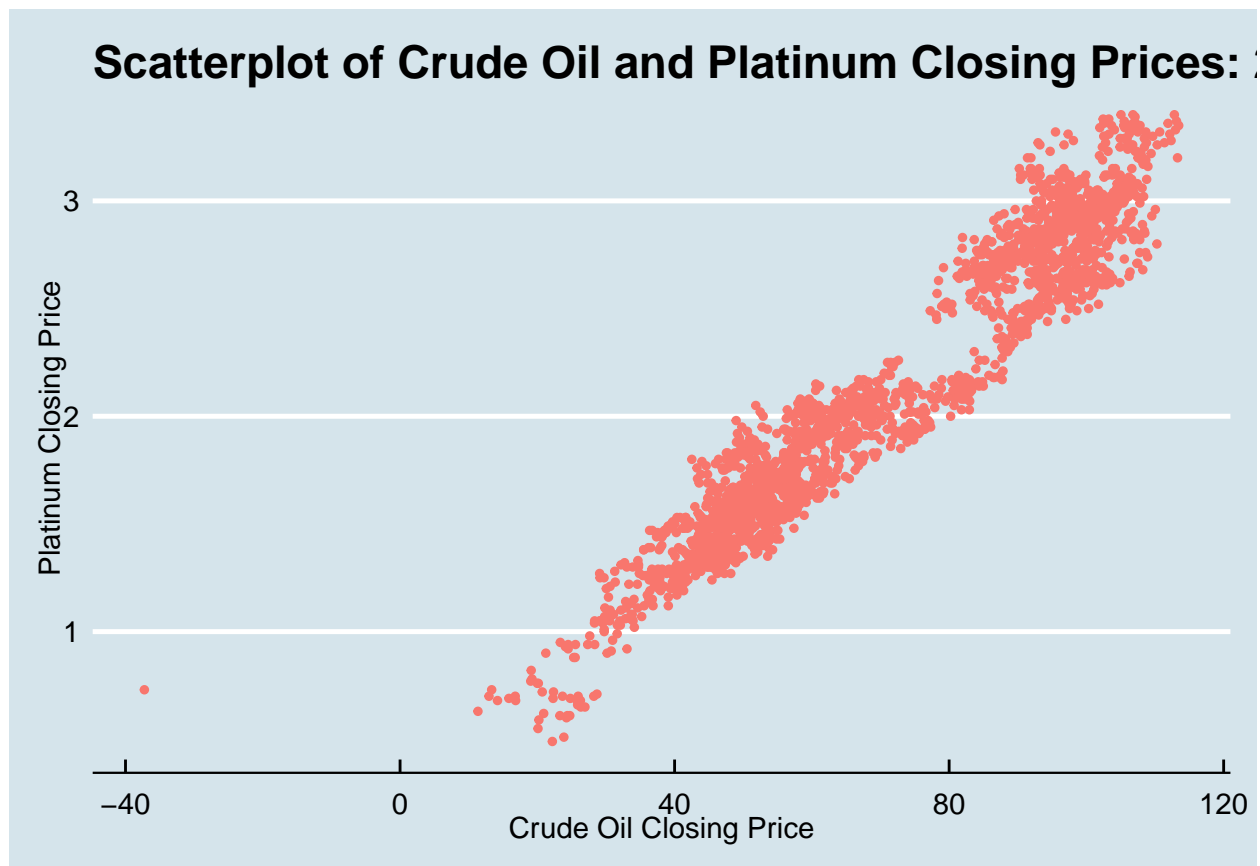


Saving 6.5 x 4.5 in image




```
## Saving 6.5 x 4.5 in image
```

```
crude_oil_platinum_scatterplot
```




```
ggsave("fig/scatterplot_crude_oil_versus_platinum.png")
```


```
## Saving 6.5 x 4.5 in image
```



Saving 6.5 x 4.5 in image




Saving 6.5 x 4.5 in image



```
## Saving 6.5 x 4.5 in image
```

```
crude_oil_soybeans_scatterplot
```

Scatterplot of Crude Oil and Soybean Closing Prices: 2



```
ggsave("fig/scatterplot_crude_oil_versus_soybean.png")
```

```
## Saving 6.5 x 4.5 in image
```

```
## Warning: The 'i' argument of '[' can't be a matrix as of tibble 3.0.0.
```

```
## Convert to a vector.
```

```
## This warning is displayed once every 8 hours.
```

```
## Call 'lifecycle::last_warnings()' to see where this warning was generated.
```

```
### This section evaluates a set of fifteen potentially viable algorithms for predicting crude oil closing prices
```

```
### Each of the first nine algorithms considers crude oil in isolation from other globally traded commodities
```

```
### In isolation, the available predictors include both the cyclical and the linear components of time
```

```
### One cyclical component is each observation's day of the trading week (Monday, Tuesday,...Friday)
```

```
### Others are the month (1,2,..12) and the quarter (1,2,..4) of each year on the 2010-2020 timeline
```

```
### Linear components are the year (e.g., 2010) and month (e.g., 8/2010) of each observation
```

```
#### ALGORITHM 01 (AVERAGE)
```

```
##### This algorithm predicts crude oil closing prices based solely on the average crude oil price for the training set
```

```
##### It will serve as a baseline from which to compare the analytical viability of other, more sophisticated algorithms
```

```
mu <- mean(crude_oil_train$closing_price)
```

```
predicted_price_algorithm_01 <- mu
```

```
RMSE01 <- RMSE(predicted_price_algorithm_01, crude_oil_test$closing_price)
```

```
RMSE01
```

```
## [1] 23.51292
```

```
#### ALGORITHM 02 (WEEKDAY EFFECTS)
```

```
##### This algorithm predicts crude oil closing prices based on average crude oil price plus the effect
```

```
crude_oil_average_by_day_of_the_week <- crude_oil_train %>%  
  group_by(date_weekday) %>%  
  summarize(b_w = mean(closing_price))
```

```
## 'summarise()' ungrouping output (override with '.groups' argument)
```

```
predicted_price_algorithm_02 <- crude_oil_test %>%  
  left_join(crude_oil_average_by_day_of_the_week, by= 'date_weekday') %>%  
  mutate(pred = b_w) %>%  
  pull(pred)
```

```
RMSE02 <- RMSE(predicted_price_algorithm_02, crude_oil_test$closing_price)
```

```
RMSE02
```

```
## [1] 23.51689
```

```
#### ALGORITHM 03 (MONTH OF THE YEAR EFFECTS)
```

```
##### This algorithm predicts crude oil closing prices based on average crude oil price plus the effect
```

```
crude_oil_average_by_month <- crude_oil_train %>%  
  group_by(date_month_of_the_year) %>%  
  summarize(b_m_y = mean(closing_price))
```

```
## 'summarise()' ungrouping output (override with '.groups' argument)
```

```
predicted_price_algorithm_03 <- crude_oil_test %>%  
  left_join(crude_oil_average_by_month, by= 'date_month_of_the_year') %>%  
  mutate(pred = b_m_y) %>%  
  pull(pred)
```

```
RMSE03 <- RMSE(predicted_price_algorithm_03, crude_oil_test$closing_price)
```

```
RMSE03
```

```
## [1] 23.54449
```

```
#### ALGORITHM 04 (QUARTER OF THE YEAR EFFECTS)
##### This algorithm predicts crude oil closing prices based on average crude oil price plus the effect.
```

```
crude_oil_average_by_quarter_of_the_year <- crude_oil_train %>%
  group_by(date_quarter_of_the_year) %>%
  summarize(b_q_y = mean(closing_price))
```

```
## 'summarise()' ungrouping output (override with '.groups' argument)
```

```
predicted_price_algorithm_04 <- crude_oil_test %>%
  left_join(crude_oil_average_by_quarter_of_the_year, by= 'date_quarter_of_the_year') %>%
  mutate(pred = b_q_y) %>%
  pull(pred)
```

```
RMSE04 <- RMSE(predicted_price_algorithm_04, crude_oil_test$closing_price)
```

```
RMSE04
```

```
## [1] 23.51713
```

```
#### ALGORITHM 05 (YEAR EFFECTS)
##### This algorithm predicts crude oil closing prices based on average crude oil price plus the effect.
```

```
crude_oil_average_by_year <- crude_oil_train %>%
  group_by(date_year) %>%
  summarize(b_y = mean(closing_price))
```

```
## 'summarise()' ungrouping output (override with '.groups' argument)
```

```
predicted_price_algorithm_05 <- crude_oil_test %>%
  left_join(crude_oil_average_by_year, by= 'date_year') %>%
  pull(b_y)
```

```
RMSE05 <- RMSE(predicted_price_algorithm_05, crude_oil_test$closing_price)
```

```
RMSE05
```

```
## [1] 8.268093
```

```
#### ALGORITHM 06 (MONTH EFFECTS)
##### This algorithm predicts crude oil closing prices based on average crude oil price plus the effect.
```

```
crude_oil_average_by_month <- crude_oil_train %>%
  group_by(date_month) %>%
  summarize(b_m = mean(closing_price))
```

```
## 'summarise()' ungrouping output (override with '.groups' argument)
```



```

predicted_price_algorithm_06 <- crude_oil_test %>%
  left_join(crude_oil_average_by_month, by= 'date_month') %>%
  mutate(pred = b_m) %>%
  pull(pred)

RMSE06 <- RMSE(predicted_price_algorithm_06, crude_oil_test$closing_price)

RMSE06

```

```
## [1] 3.800524
```

```


#### ALGORITHM 07 (RANDOM FOREST - TIME)
##### This algorithm predicts crude oil closing prices based on a Random Forest of time components
##### The first step is to look at a Random Forest incorporating all of above-listed components of ti

nodesize <- seq(1, 5, 1)

rmses <- sapply(nodesize, function(n){
  rf_time = randomForest(closing_price ~
                        date_weekday +
                        date_month_of_the_year +
                        date_quarter_of_the_year +
                        date_year +
                        date_month,
                        data = crude_oil_train, nodesize = n)
  pred <- predict(rf_time, newdata = crude_oil_train)
  RMSE(pred, crude_oil_train$closing_price)
})

qplot(nodesize, rmses)

```



```
nodesize[which.min(rmse)]
```


```
## [1] 5
```

```
rf_time = randomForest(closing_price ~
                        date_weekday +
                        date_month_of_the_year +
                        date_quarter_of_the_year +
                        date_year +
                        date_month,
                        data = crude_oil_train, nodesize = nodesize[which.min(rmse)])
```

The second is to build a plot to help determine which of those components is/are the most important

```
varImpPlot(rf_time)
```

rf_time




Based on that plot, Random Forest will be revised to focus on the month and year

```
nodesize <- seq(1, 2, 1)

rmsees <- sapply(nodesize, function(n){
  rf_time = randomForest(closing_price ~
                        date_month +
                        date_year,
                        data = crude_oil_train, nodesize = n)
  pred <- predict(rf_time, newdata = crude_oil_train)
  RMSE(pred, crude_oil_train$closing_price)
})

qplot(nodesize, rmsees)
```



```
nodesize[which.min(rmses)]
```

```
## [1] 2
```

```
rf_time = randomForest(closing_price ~
                        date_month +
                        date_year,
                        data = crude_oil_train, nodesize = nodesize[which.min(rmses)])

pred <- predict(rf_time, newdata = crude_oil_train)

RMSE(pred, crude_oil_train$closing_price)
```

```
## [1] 3.850352
```

```
predicted_price_algorithm_07 <- predict(rf_time, newdata = crude_oil_test)

RMSE07 <- RMSE(predicted_price_algorithm_07, crude_oil_test$closing_price)

RMSE07
```


```
## [1] 4.716089
```

```
#### ALGORITHM 08 (KNN - TIME)
```

```
##### This algorithm predicts crude oil closing prices based on a K-Nearest Neighbors (KNN) model of th
```

```
knn_time <- train(closing_price ~  
  date_year +  
  date_month,  
  method = "knn",  
  data = crude_oil_train)
```

```
ggplot(knn_time, highlight = TRUE)
```



```
pred <- predict(knn_time, newdata = crude_oil_train)
```

```
RMSE(pred, crude_oil_train$closing_price)
```

```
## [1] 2.572097
```

```
predicted_price_algorithm_08 <- predict(knn_time, newdata = crude_oil_test)
```

```
RMSE08 <- RMSE(predicted_price_algorithm_08, crude_oil_test$closing_price)
```

```
RMSE08
```

```
## [1] 3.768831
```

```
#### ALGORITHM 09 (RANGER - TIME)
##### This algorithm predicts crude oil closing prices based on a Ranger model of the two key time comp

ranger_time <- train(closing_price ~
                     date_year +
                     date_month,
                     method = "ranger",
                     data = crude_oil_train)
```

note: only 1 unique complexity parameters in default grid. Truncating the grid to 1 .

```
pred <- predict(ranger_time, newdata = crude_oil_train)

RMSE(pred, crude_oil_train$closing_price)
```

```
## [1] 2.572239
```

```
predicted_price_algorithm_09 <- predict(ranger_time, newdata = crude_oil_test)

RMSE09 <- RMSE(predicted_price_algorithm_09, crude_oil_test$closing_price)

RMSE09
```

```
## [1] 3.772397
```


Each of the next three algorithms considers crude oil in comparison to other commodities from the e
 ### The available predictors are the key components of time and the closing prices of natural gas, heat

```
#### ALGORITHM 10 (RANDOM FOREST - TIME AND ENERGY)
##### This algorithm predicts crude oil closing prices based on a Random Forest of time components and
##### The first step is to look at a random forest incorporating year, month, and closing prices of n
```

```
nodesize <- seq(1, 5, 1)

rmsees <- sapply(nodesize, function(n){
  rf_time_energy = randomForest(closing_price ~
                                date_year +
                                date_month +
                                natural_gas_closing_price +
                                heating_oil_closing_price +
                                gasoline_closing_price,
                                data = crude_oil_train, nodesize = n)
  pred <- predict(rf_time_energy, newdata = crude_oil_train)
  RMSE(pred, crude_oil_train$closing_price)
})

qplot(nodesize, rmsees)
```



```
nodesize[which.min(rmse)]
```

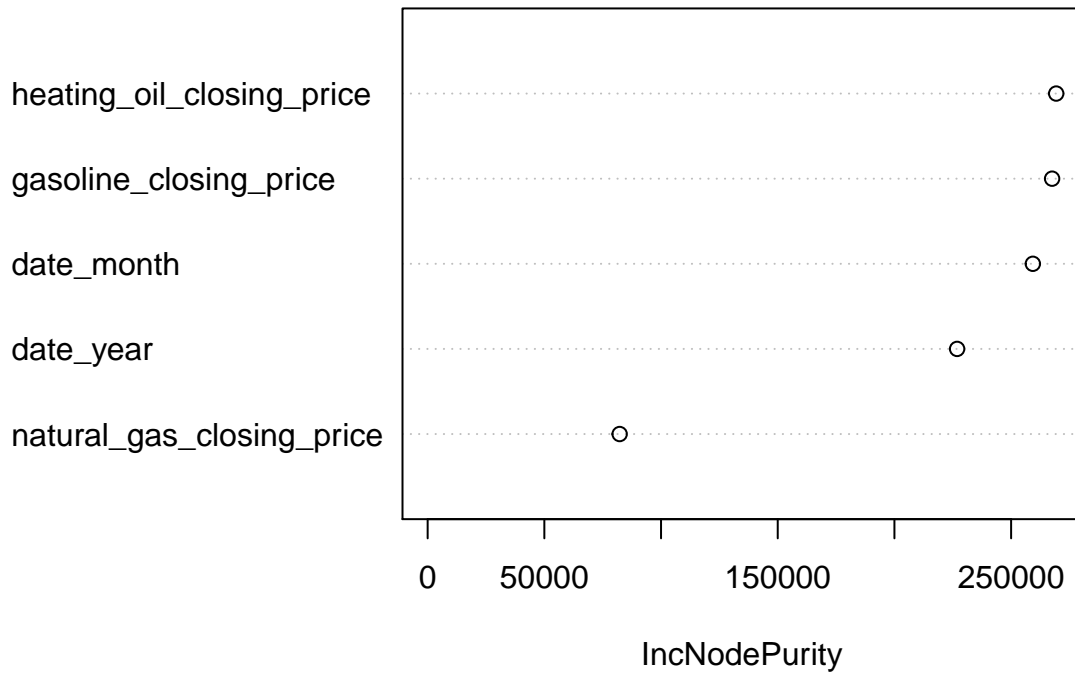
```
## [1] 2
```

```
rf_time_energy = randomForest(closing_price ~
                              date_year +
                              date_month +
                              natural_gas_closing_price +
                              heating_oil_closing_price +
                              gasoline_closing_price,
                              data = crude_oil_train, nodesize = nodesize[which.min(rmse)])
```

The second is to build a plot to help determine which of those components is/are the most important

```
varImpPlot(rf_time_energy)
```

rf_time_energy




Based on that plot, Random Forest will be revised to focus on year, month, heating oil, and gas

```
nodesize <- seq(1, 4, 1)

rmsees <- sapply(nodesize, function(n){
  rf_time_energy = randomForest(closing_price ~
                                date_year +
                                date_month +
                                heating_oil_closing_price +
                                gasoline_closing_price,
                                data = crude_oil_train, nodesize = n)
  pred <- predict(rf_time_energy, newdata = crude_oil_train)
  RMSE(pred, crude_oil_train$closing_price)
})

qplot(nodesize, rmsees)
```

```
nodesize[which.min(rmses)]
```

```
## [1] 1
```

```
rf_time_energy = randomForest(closing_price ~
                              date_year +
                              date_month +
                              heating_oil_closing_price +
                              gasoline_closing_price,
                              data = crude_oil_train, nodesize = nodesize[which.min(rmses)])

pred <- predict(rf_time_energy, newdata = crude_oil_train)

RMSE(pred, crude_oil_train$closing_price)
```

```
## [1] 1.55598
```

```
predicted_price_algorithm_10 <- predict(rf_time_energy, newdata = crude_oil_test)

RMSE10 <- RMSE(predicted_price_algorithm_10, crude_oil_test$closing_price)


RMSE10
```

```
## [1] 3.297452
```

```
#### ALGORITHM 11 (KNN - TIME AND ENERGY)
##### This algorithm predicts crude oil closing prices based on a KNN model of the four key time and en

knn_time_energy <- train(closing_price ~
  date_year +
  date_month +
  heating_oil_closing_price +
  gasoline_closing_price,
  method = "knn",
  data = crude_oil_train)

ggplot(knn_time_energy, highlight = TRUE)
```



```
pred <- predict(knn_time_energy, newdata = crude_oil_train)
RMSE(pred, crude_oil_train$closing_price)
```

```
## [1] 1.337656
```

```
predicted_price_algorithm_11 <- predict(knn_time_energy, newdata = crude_oil_test)
RMSE11 <- RMSE(predicted_price_algorithm_11, crude_oil_test$closing_price)
RMSE11
```

```
## [1] 2.956255
```

```
#### ALGORITHM 12 (RANGER - TIME AND ENERGY)
##### This algorithm predicts crude oil closing prices based on a RANGER model of the key time and energy
```

```
ranger_time_energy<- train(closing_price ~
                           date +
                           date_year +
                           heating_oil_closing_price +
                           gasoline_closing_price,
                           method = "ranger",
                           data = crude_oil_train)

pred <- predict(ranger_time_energy, newdata = crude_oil_train)

RMSE(pred, crude_oil_train$closing_price)
```

```
## [1] 0.672647
```

```
predicted_price_algorithm_12 <- predict(ranger_time_energy, newdata = crude_oil_test)

RMSE12 <- RMSE(predicted_price_algorithm_12, crude_oil_test$closing_price)

RMSE12
```

```
## [1] 2.963854
```


```
### Each of the final three algorithms considers crude oil in comparison to other commodities from the
### The predictors include not only time but the closing prices of natural gas, heating oil, gasoline,
```

```
#### ALGORITHM 13 (RANDOM FOREST - TIME, ENERGY, PRECIOUS METALS, AGRICULTURE)
##### This algorithm predicts crude oil closing prices based on a Random Forest of time, energy, precious metals, and agriculture
##### The first step is to look at a random forest incorporating date, year, and closing prices of energy, precious metals, and agriculture
```

```
nodesize <- seq(1, 10, 1)

rmsees <- sapply(nodesize, function(n){
  rf_time_energy_metals_agriculture = randomForest(closing_price ~
                                                    date_year +
                                                    date_month +
                                                    heating_oil_closing_price +
                                                    gasoline_closing_price +
                                                    gold_closing_price +
                                                    silver_closing_price +
                                                    platinum_closing_price +
                                                    wheat_closing_price +
                                                    rice_closing_price +
                                                    soybeans_closing_price,
                                                    data = crude_oil_train, nodesize = n)
  pred <- predict(rf_time_energy_metals_agriculture, newdata = crude_oil_train)
  RMSE(pred, crude_oil_train$closing_price)
})
```

```
qplot(nodesize, rmses)
```



```
nodesize[which.min(rmses)]
```

```
## [1] 1
```


```
rf_time_energy_metals_agriculture = randomForest(closing_price ~
  date_year +
  date_month +
  heating_oil_closing_price +
  gasoline_closing_price +
  gold_closing_price +
  silver_closing_price +
  platinum_closing_price +
  wheat_closing_price +
  rice_closing_price +
  soybeans_closing_price,
  data = crude_oil_train, nodesize = nodesize[which.min(rmses)])

##### The second is to build a plot to help determine which of those components is/are the most important

varImpPlot(rf_time_energy_metals_agriculture)
```

rf_time_energy_metals_agriculture

date_month
heating_oil_closing_price
platinum_closing_price
date_year
gasoline_closing_price
soybeans_closing_price
wheat_closing_price
silver_closing_price
rice_closing_price
gold_closing_price




Based on that plot, Random Forest will be revised to focus on year, month, heating oil, gasoline

```
nodesize <- seq(1, 6, 1)

rmsees <- sapply(nodesize, function(n){
  rf_time_energy_metals_agriculture = randomForest(closing_price ~
                                                    date_year +
                                                    date_month +
                                                    heating_oil_closing_price +
                                                    gasoline_closing_price +
                                                    platinum_closing_price +
                                                    soybeans_closing_price,
                                                    data = crude_oil_train, nodesize = n)
  pred <- predict(rf_time_energy_metals_agriculture, newdata = crude_oil_train)
  RMSE(pred, crude_oil_train$closing_price)
})

qplot(nodesize, rmsees)
```



```
nodesize[which.min(rmses)]
```

```
## [1] 1
```

```
rf_time_energy_metals_agriculture = randomForest(closing_price ~
  date_year +
  date_month +
  heating_oil_closing_price +
  gasoline_closing_price +
  platinum_closing_price +
  soybeans_closing_price,
  data = crude_oil_train, nodesize = nodesize[which.min(rmses)])

pred <- predict(rf_time_energy_metals_agriculture, newdata = crude_oil_train)

RMSE(pred, crude_oil_train$closing_price)
```

```
## [1] 0.6148404
```

```
predicted_price_algorithm_13 <- predict(rf_time_energy_metals_agriculture, newdata = crude_oil_test)

RMSE13 <- RMSE(predicted_price_algorithm_13, crude_oil_test$closing_price)

RMSE13
```


```
## [1] 2.951125
```

```
#### ALGORITHM 14 (KNN - TIME, ENERGY, PRECIOUS METALS, AGRICULTURE)
```

```
##### This algorithm predicts crude oil closing prices based on a KNN model incorporating date, heating
```

```
knn_time_energy_metals_agriculture <- train(closing_price ~  
      date_year +  
      date_month +  
      heating_oil_closing_price +  
      gasoline_closing_price +  
      platinum_closing_price +  
      soybeans_closing_price,  
      method = "knn",  
      data = crude_oil_train)
```

```
ggplot(knn_time_energy_metals_agriculture, highlight = TRUE)
```



```
pred <- predict(knn_time_energy_metals_agriculture, newdata = crude_oil_train)
```

```
RMSE(pred, crude_oil_train$closing_price)
```

```
## [1] 1.689948
```

```

predicted_price_algorithm_14 <- predict(knn_time_energy_metals_agriculture, newdata = crude_oil_test)

RMSE14 <- RMSE(predicted_price_algorithm_14, crude_oil_test$closing_price)

RMSE14

```

```
## [1] 3.298945
```

```

#### ALGORITHM 15 (RANGER - TIME, ENERGY, PRECIOUS METALS, AGRICULTURE)
##### This algorithm predicts crude oil closing prices based on a RANGER model incorporating date, heat

```

```

ranger_time_energy_metals_agriculture <- train(closing_price ~
      date_year +
      date_month +
      heating_oil_closing_price +
      gasoline_closing_price +
      platinum_closing_price +
      soybeans_closing_price,
      method = "ranger",
      data = crude_oil_train)

pred <- predict(ranger_time_energy_metals_agriculture, newdata = crude_oil_train)

RMSE(pred, crude_oil_train$closing_price)

```

```
## [1] 0.8545436
```

```

predicted_price_algorithm_15 <- predict(ranger_time_energy_metals_agriculture, newdata = crude_oil_test)

RMSE15 <- RMSE(predicted_price_algorithm_15, crude_oil_test$closing_price)

RMSE15

```

```
## [1] 2.930749
```

```
### SUMMARY OF FINDINGS
```

```

ALGORITHM_tab <- c("ALGORITHM 01 - Average Closing Price",
  "ALGORITHM 02 - Average Closing Price plus Day of the Week Effects",
  "ALGORITHM 03 - Average Closing Price plus Month of the Year Effects",
  "ALGORITHM 04 - Average Closing Price plus Quarter of the Year Effects",
  "ALGORITHM 05 - Average Closing Price plus Year Effects",
  "ALGORITHM 06 - Average Closing Price Plus Month Effects",
  "ALGORITHM 07 - Random Forest of Time Components",
  "ALGORITHM 08 - KNN Model of Time Components",
  "ALGORITHM 09 - Ranger Model of Time Components",
  "ALGORITHM 10 - Random Forest of Time and Energy Components",
  "ALGORITHM 11 - KNN Model of Time and Energy Components",
  "ALGORITHM 12 - Ranger Model of Time and Energy Components",
  "ALGORITHM 13 - Random Forest of Time Components and Energy, Precious Metals, and Agr",
  "ALGORITHM 14 - KNN Model of Time Components and Energy, Precious Metals, and Agricul

```



```

"ALGORITHM 15 - Ranger Model of Time Components and Energy, Precious Metals,and Agri

RMSE_tab <- c(RMSE01,
             RMSE02,
             RMSE03,
             RMSE04,
             RMSE05,
             RMSE06,
             RMSE07,
             RMSE08,
             RMSE09,
             RMSE10,
             RMSE11,
             RMSE12,
             RMSE13,
             RMSE14,
             RMSE15)

Results <- data.frame(ALGORITHM_tab, RMSE_tab)

Results %>%
  arrange(RMSE_tab) %>%
  print(right = FALSE)

```

```

##   ALGORITHM_tab
## 1  ALGORITHM 15 - Ranger Model of Time Components and Energy, Precious Metals,and Agriculture Commod
## 2  ALGORITHM 13 - Random Forest of Time Components and Energy, Precious Metals,and Agriculture Commo
## 3  ALGORITHM 11 - KNN Model of Time and Energy Components
## 4  ALGORITHM 12 - Ranger Model of Time and Energy Components
## 5  ALGORITHM 10 - Random Forest of Time and Energy Components
## 6  ALGORITHM 14 - KNN Model of Time Components and Energy, Precious Metals,and Agriculture Commoditi
## 7  ALGORITHM 08 - KNN Model of Time Components
## 8  ALGORITHM 09 - Ranger Model of Time Components
## 9  ALGORITHM 06 - Average Closing Price Plus Month Effects
## 10 ALGORITHM 07 - Random Forest of Time Components
## 11 ALGORITHM 05 - Average Closing Price plus Year Effects
## 12 ALGORITHM 01 - Average Closing Price
## 13 ALGORITHM 02 - Average Closing Price plus Day of the Week Effects
## 14 ALGORITHM 04 - Average Closing Price plus Quarter of the Year Effects
## 15 ALGORITHM 03 - Average Closing Price plus Month of the Year Effects
##   RMSE_tab
## 1   2.930749
## 2   2.951125
## 3   2.956255
## 4   2.963854
## 5   3.297452
## 6   3.298945
## 7   3.768831
## 8   3.772397
## 9   3.800524
## 10  4.716089
## 11  8.268093
## 12 23.512925

```

```
## 13 23.516889
## 14 23.517132
## 15 23.544488
```